

▼ Loading the data

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from google.colab import files
from datetime import datetime
import seaborn as sns
%matplotlib inline

import io

# Load the data
local_file = files.upload()
train_data = io.BytesIO(local_file['train1.csv'])
train_data2 = io.BytesIO(local_file['train.csv'])
df1 = pd.read_csv(train_data)
df2 = pd.read_csv(train_data2)
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving train.csv to train.csv
Saving train1.csv to train1.csv

▼ Data integration

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 367 entries, 0 to 366
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                367 non-null   object
1   Gender                 356 non-null   object
2   Married                367 non-null   object
3   Dependents             357 non-null   object
4   Education              367 non-null   object
5   Self_Employed          344 non-null   object
6   ApplicantIncome        367 non-null   int64
7   CoapplicantIncome      367 non-null   int64
```

```

8   LoanAmount      362 non-null    float64
9   Loan_Amount_Term  361 non-null    float64
10  Credit_History    338 non-null    float64
11  Property_Area     367 non-null    object
dtypes: float64(3), int64(2), object(7)
memory usage: 34.5+ KB

```

```
df2.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                614 non-null    object
1   Gender                 601 non-null    object
2   Married                611 non-null    object
3   Dependents             599 non-null    object
4   Education              614 non-null    object
5   Self_Employed          582 non-null    object
6   ApplicantIncome        614 non-null    int64
7   CoapplicantIncome      614 non-null    float64
8   LoanAmount             592 non-null    float64
9   Loan_Amount_Term       600 non-null    float64
10  Credit_History         564 non-null    float64
11  Property_Area          614 non-null    object
12  Loan_Status            614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB

```

```
frames = [df1, df2]
```

```
df = pd.concat(frames)
df
```

| | Loan_ID | Gender | Married | Dependents | Education | Self_Employed | ApplicantIncome |
|---|----------|--------|---------|------------|-----------|---------------|-----------------|
| 0 | LP001015 | Male | Yes | 0 | Graduate | No | 5720 |
| 1 | LP001022 | Male | Yes | 1 | Graduate | No | 3076 |
| 2 | LP001031 | Male | Yes | 2 | Graduate | No | 5000 |

```
df.tail()
```

| | Loan_ID | Gender | Married | Dependents | Education | Self_Employed | ApplicantIncome |
|-----|----------|--------|---------|------------|-----------|---------------|-----------------|
| 609 | LP002978 | Female | No | 0 | Graduate | No | 2900 |
| 610 | LP002979 | Male | Yes | 3+ | Graduate | No | 4106 |
| 611 | LP002983 | Male | Yes | 1 | Graduate | No | 8072 |
| 612 | LP002984 | Male | Yes | 2 | Graduate | No | 7583 |
| 613 | LP002990 | Female | No | 0 | Graduate | Yes | 4583 |

▼ Data cleaning

```
df.isnull().sum()
```

```
Loan_ID      0
Gender       24
Married       3
Dependents   25
Education     0
Self_Employed 55
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount   27
Loan_Amount_Term 20
Credit_History 79
Property_Area 0
Loan_Status  367
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 981 entries, 0 to 613
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Loan_ID             981 non-null   object
1   Gender              957 non-null   object
2   Married             978 non-null   object
```

```

3   Dependents      956 non-null    object
4   Education       981 non-null    object
5   Self_Employed   926 non-null    object
6   ApplicantIncome 981 non-null    int64
7   CoapplicantIncome 981 non-null    float64
8   LoanAmount      954 non-null    float64
9   Loan_Amount_Term 961 non-null    float64
10  Credit_History   902 non-null    float64
11  Property_Area    981 non-null    object
12  Loan_Status      614 non-null    object
dtypes: float64(4), int64(1), object(8)
memory usage: 107.3+ KB

```

Converting the string values to numeric values.

```

numeric_gender = {'Female': 1, 'Male': 2}
df ['Gender'] = df['Gender'].map(numeric_gender)
numeric_married = {'Yes': 1, 'No': 2}
df ['Married'] = df['Married'].map(numeric_married)
numeric_edu = {'Graduate': 1, 'Not Graduate': 2}
df ['Education'] = df['Education'].map(numeric_edu)
numeric_self = {'Yes': 1, 'No': 2}
df ['Self_Employed'] = df['Self_Employed'].map(numeric_self)
numeric_loan = {'Y': 1, 'N': 2}
df ['Loan_Status'] = df['Loan_Status'].map(numeric_loan)
numeric_property = {'Rural': 1, 'Urban': 2, 'Semiurban': 3}
df ['Property_Area'] = df['Property_Area'].map(numeric_property)
#numeric_d = {'3+': 3}
#df ['Dependents'] = df['Dependents'].map(numeric_d)

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 981 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               981 non-null    object
1   Gender                957 non-null    float64
2   Married               978 non-null    float64
3   Dependents            956 non-null    object
4   Education             981 non-null    int64
5   Self_Employed         926 non-null    float64
6   ApplicantIncome       981 non-null    int64
7   CoapplicantIncome     981 non-null    float64
8   LoanAmount            954 non-null    float64
9   Loan_Amount_Term      961 non-null    float64
10  Credit_History        902 non-null    float64
11  Property_Area         981 non-null    int64
12  Loan_Status           614 non-null    float64
dtypes: float64(8), int64(3), object(2)
memory usage: 107.3+ KB

```

```
df.isnull().sum()
```

```
Loan_ID      0
Gender       24
Married       3
Dependents   25
Education     0
Self_Employed 55
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount   27
Loan_Amount_Term 20
Credit_History 79
Property_Area 0
Loan_Status  367
dtype: int64
```

▼ Data analysis:

```
df.describe()
```

| | Gender | Married | Education | Self_Employed | ApplicantIncome | CoapplicantIncome |
|--------------|------------|------------|------------|---------------|-----------------|-------------------|
| count | 957.000000 | 978.000000 | 981.000000 | 926.000000 | 981.000000 | 981.000000 |
| mean | 1.809822 | 1.354806 | 1.222222 | 1.871490 | 5179.795107 | 1601.916000 |
| std | 0.392646 | 0.478699 | 0.415952 | 0.334837 | 5695.104533 | 2718.772000 |
| min | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 0.000000 | 0.000000 |
| 25% | 2.000000 | 1.000000 | 1.000000 | 2.000000 | 2875.000000 | 0.000000 |
| 50% | 2.000000 | 1.000000 | 1.000000 | 2.000000 | 3800.000000 | 1110.000000 |
| 75% | 2.000000 | 2.000000 | 1.000000 | 2.000000 | 5516.000000 | 2365.000000 |
| max | 2.000000 | 2.000000 | 2.000000 | 2.000000 | 81000.000000 | 41667.000000 |

Here we can see the shape of our data with the `.shape`. Here we see (614, 13) this means that we have a 614 rows and 13 columns

```
df.shape
```

```
(981, 13)
```

Here we can see the shape of our test data with the .shape. Here we see (367, 12) this means that

To view what data that is stored we can use .columns. This will return the columns of our data

```
df.columns
```

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
       'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

To look at the data we'll use the .head() method from pandas. This will show us the first 5 items in our dataframe.

```
df.head()
```

| | Loan_ID | Gender | Married | Dependents | Education | Self_Employed | ApplicantIncome | Co |
|----------|----------|--------|---------|------------|-----------|---------------|-----------------|----|
| 0 | LP001015 | 2.0 | 1.0 | 0 | 1 | 2.0 | 5720 | |
| 1 | LP001022 | 2.0 | 1.0 | 1 | 1 | 2.0 | 3076 | |
| 2 | LP001031 | 2.0 | 1.0 | 2 | 1 | 2.0 | 5000 | |
| 3 | LP001035 | 2.0 | 1.0 | 2 | 1 | 2.0 | 2340 | |
| 4 | LP001051 | 2.0 | 2.0 | 0 | 2 | 2.0 | 3276 | |

```
df.tail()
```

| | Loan_ID | Gender | Married | Dependents | Education | Self_Employed | ApplicantIncome | Co |
|------------|----------|--------|---------|------------|-----------|---------------|-----------------|----|
| 609 | LP002978 | 1.0 | 2.0 | 0 | 1 | 2.0 | 2900 | |
| 610 | LP002979 | 2.0 | 1.0 | 3+ | 1 | 2.0 | 4106 | |
| 611 | LP002983 | 2.0 | 1.0 | 1 | 1 | 2.0 | 8072 | |
| 612 | LP002984 | 2.0 | 1.0 | 2 | 1 | 2.0 | 7583 | |
| 613 | LP002990 | 1.0 | 2.0 | 0 | 1 | 1.0 | 4583 | |

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 981 entries, 0 to 613
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
#   ...
```

```
-----
0   Loan_ID          981 non-null    object
1   Gender           957 non-null    float64
2   Married          978 non-null    float64
3   Dependents       956 non-null    object
4   Education         981 non-null    int64
5   Self_Employed    926 non-null    float64
6   ApplicantIncome  981 non-null    int64
7   CoapplicantIncome 981 non-null    float64
8   LoanAmount       954 non-null    float64
9   Loan_Amount_Term 961 non-null    float64
10  Credit_History   902 non-null    float64
11  Property_Area    981 non-null    int64
12  Loan_Status      614 non-null    float64
dtypes: float64(8), int64(3), object(2)
memory usage: 107.3+ KB
```

```
df.sort_values('Loan_Status', ascending = True)[:50]
```

| | Loan_ID | Gender | Married | Dependents | Education | Self_Employed | ApplicantIncome |
|------------|----------|--------|---------|------------|-----------|---------------|-----------------|
| 0 | LP001002 | 2.0 | 2.0 | 0 | 1 | 2.0 | 5849 |
| 380 | LP002226 | 2.0 | 1.0 | 0 | 1 | NaN | 3333 |
| 379 | LP002225 | 2.0 | 1.0 | 2 | 1 | 2.0 | 5391 |
| 377 | LP002223 | 2.0 | 1.0 | 0 | 1 | 2.0 | 4310 |
| 376 | LP002219 | 2.0 | 1.0 | 3+ | 1 | 2.0 | 8750 |
| 375 | LP002211 | 2.0 | 1.0 | 0 | 1 | 2.0 | 4817 |
| 374 | LP002209 | 1.0 | 2.0 | 0 | 1 | NaN | 2764 |
| 372 | LP002201 | 2.0 | 1.0 | 2 | 1 | 1.0 | 9323 |
| 371 | LP002197 | 2.0 | 1.0 | 2 | 1 | 2.0 | 5185 |
| 370 | LP002194 | 1.0 | 2.0 | 0 | 1 | 1.0 | 15759 |
| 368 | LP002190 | 2.0 | 1.0 | 1 | 1 | 2.0 | 6325 |

```
df.sort_values('Education')[:50]
```


| | Loan_ID | Gender | Married | Dependents | Education | Self_Employed | ApplicantIncome |
|-----|----------|--------|---------|------------|-----------|---------------|-----------------|
| 0 | LP001015 | 2.0 | 1.0 | 0 | 1 | 2.0 | 5720 |
| 234 | LP001778 | 2.0 | 1.0 | 1 | 1 | 2.0 | 3155 |
| 235 | LP001784 | 2.0 | 1.0 | 1 | 1 | 2.0 | 5500 |
| 236 | LP001786 | 2.0 | 1.0 | 0 | 1 | NaN | 5746 |
| 237 | LP001788 | 1.0 | 2.0 | 0 | 1 | 1.0 | 3463 |
| 238 | LP001790 | 1.0 | 2.0 | 1 | 1 | 2.0 | 3812 |
| 239 | LP001792 | 2.0 | 1.0 | 1 | 1 | 2.0 | 3315 |
| 240 | LP001798 | 2.0 | 1.0 | 2 | 1 | 2.0 | 5819 |
| 242 | LP001806 | 2.0 | 2.0 | 0 | 1 | 2.0 | 2965 |
| 243 | LP001807 | 2.0 | 1.0 | 2 | 1 | 1.0 | 6250 |
| 245 | LP001813 | 2.0 | 2.0 | 0 | 1 | 1.0 | 6050 |
| 233 | LP001776 | 1.0 | 2.0 | 0 | 1 | 2.0 | 8333 |
| 246 | LP001814 | 2.0 | 1.0 | 2 | 1 | 2.0 | 9703 |
| 249 | LP001825 | 2.0 | 1.0 | 0 | 1 | 2.0 | 1809 |
| 251 | LP001836 | 1.0 | 2.0 | 2 | 1 | 2.0 | 3427 |
| 254 | LP001844 | 2.0 | 2.0 | 0 | 1 | 1.0 | 16250 |
| 255 | LP001846 | 1.0 | 2.0 | 3+ | 1 | 2.0 | 3083 |
| 257 | LP001854 | 2.0 | 1.0 | 3+ | 1 | 2.0 | 5250 |
| 258 | LP001859 | 2.0 | 1.0 | 0 | 1 | 2.0 | 14683 |
| 260 | LP001865 | 2.0 | 1.0 | 1 | 1 | 2.0 | 6083 |
| 261 | LP001868 | 2.0 | 2.0 | 0 | 1 | 2.0 | 2060 |
| 262 | LP001870 | 1.0 | 2.0 | 1 | 1 | 2.0 | 3481 |
| 263 | LP001871 | 1.0 | 2.0 | 0 | 1 | 2.0 | 7200 |
| 248 | LP001824 | 2.0 | 1.0 | 1 | 1 | 2.0 | 2882 |
| 231 | LP001768 | 2.0 | 1.0 | 0 | 1 | NaN | 3716 |
| 230 | LP001765 | 2.0 | 1.0 | 1 | 1 | 2.0 | 2491 |
| 229 | LP001761 | 2.0 | 2.0 | 0 | 1 | 1.0 | 6400 |
| 196 | LP001666 | 2.0 | 2.0 | 0 | 1 | 2.0 | 8333 |
| 198 | LP001671 | 1.0 | 1.0 | 0 | 1 | 2.0 | 3416 |
| 199 | LP001673 | 2.0 | 2.0 | 0 | 1 | 1.0 | 11000 |

| | | | | | | | |
|------------|----------|-----|-----|---|---|-----|------|
| 201 | LP001677 | 2.0 | 2.0 | 2 | 1 | 2.0 | 4923 |
| 206 | LP001693 | 1.0 | 2.0 | 0 | 1 | 2.0 | 3244 |

Here we can see one row (one person)

```
df.iloc[0]
```

```

Loan_ID      LP001015
Gender        2
Married       1
Dependents    0
Education     1
Self_Employed 2
ApplicantIncome 5720
CoapplicantIncome 0
LoanAmount    110
Loan_Amount_Term 360
Credit_History 1
Property_Area  2
Loan_Status   NaN
Name: 0, dtype: object

```

Get the unique values and their frequency of variable. (Checking how many times the certain value occurs.)

| | | | | | | | |
|------------|----------|-----|-----|---|---|-----|------|
| 204 | LP001750 | 2.0 | 1.0 | 0 | 1 | 2.0 | 6250 |
|------------|----------|-----|-----|---|---|-----|------|

```
df['Loan_Status'].value_counts()
```

```

1.0    422
2.0    192
Name: Loan_Status, dtype: int64

```

```
df['ApplicantIncome'].value_counts()
```

```

2500    13
5000    11
3333    10
3500     9
2600     8
..
5391     1
15000    1
14999    1
7830     1
1811     1
Name: ApplicantIncome, Length: 752, dtype: int64

```

```
df['Gender'].value_counts()
```

```

2.0    775
1.0    182

```

```
Name: Gender, dtype: int64
```

```
df['Married'].value_counts()
```

```
1.0    631
```

```
2.0    347
```

```
Name: Married, dtype: int64
```

```
df['CoapplicantIncome'].value_counts()
```

```
0.0    429
```

```
2500.0    6
```

```
1666.0    5
```

```
2000.0    5
```

```
2083.0    5
```

```
...
```

```
6250.0    1
```

```
1742.0    1
```

```
189.0     1
```

```
1868.0    1
```

```
4266.0    1
```

```
Name: CoapplicantIncome, Length: 437, dtype: int64
```

```
df['Dependents'].value_counts()
```

```
0    545
```

```
1    160
```

```
2    160
```

```
3+    91
```

```
Name: Dependents, dtype: int64
```

```
df['Education'].value_counts()
```

```
1    763
```

```
2    218
```

```
Name: Education, dtype: int64
```

```
df['Self_Employed'].value_counts()
```

```
2.0    807
```

```
1.0    119
```

```
Name: Self_Employed, dtype: int64
```

```
df['Loan_Status'].unique()
```

```
array([nan, 1., 2.])
```


```
df['ApplicantIncome'].unique()
```

```
array([ 5720,  3076,  5000,  2340,  3276,  2165,  2226,  3881, 13633,
```

| | | | | | | | | |
|--------|--------|--------|--------|--------|--------|--------|--------|--------|
| 2400, | 3091, | 2185, | 4166, | 12173, | 4666, | 5667, | 4583, | 3786, |
| 9226, | 1300, | 1888, | 2083, | 3909, | 3765, | 5400, | 0, | 4363, |
| 7500, | 3772, | 2942, | 2478, | 6250, | 3268, | 2783, | 2740, | 3150, |
| 7350, | 2267, | 5833, | 3643, | 5629, | 3644, | 1750, | 6500, | 3666, |
| 4260, | 4163, | 2356, | 6792, | 8000, | 2419, | 3500, | 4116, | 5293, |
| 2750, | 4402, | 3613, | 2779, | 4720, | 2415, | 7016, | 4968, | 2101, |
| 4490, | 2917, | 4700, | 3445, | 7666, | 2458, | 3250, | 4463, | 4083, |
| 3900, | 4750, | 3583, | 3189, | 6356, | 3413, | 7950, | 3829, | 72529, |
| 4136, | 8449, | 4456, | 4635, | 3571, | 3066, | 3235, | 5058, | 3188, |
| 13518, | 4364, | 4766, | 4609, | 6260, | 3333, | 9719, | 6835, | 4452, |
| 2262, | 3901, | 2687, | 2243, | 3417, | 1596, | 4513, | 4500, | 4523, |
| 4742, | 4082, | 2922, | 4167, | 4243, | 1516, | 1762, | 724, | 3125, |
| 2333, | 3350, | 2500, | 12500, | 4667, | 3073, | 3321, | 3391, | 3343, |
| 3620, | 4000, | 4258, | 2014, | 4727, | 3089, | 6794, | 32000, | 10890, |
| 12941, | 8703, | 5900, | 3071, | 2463, | 4855, | 1599, | 4246, | 4333, |
| 5823, | 7895, | 4150, | 2964, | 5583, | 2708, | 3180, | 2268, | 1141, |
| 3042, | 3564, | 3958, | 4483, | 5225, | 3017, | 2431, | 4912, | 2918, |
| 5128, | 15312, | 4334, | 4358, | 10166, | 4521, | 9167, | 13083, | 7874, |
| 3785, | 2654, | 10000, | 4796, | 2000, | 2540, | 1900, | 8706, | 2855, |
| 3016, | 3159, | 1937, | 2613, | 4960, | 3074, | 4213, | 2038, | 2362, |
| 5333, | 5384, | 5708, | 3754, | 2914, | 2747, | 7830, | 3507, | 3747, |
| 2166, | 2896, | 5062, | 5184, | 2545, | 2553, | 3436, | 2412, | 5180, |
| 14911, | 2860, | 1173, | 7600, | 2157, | 2231, | 2274, | 6166, | 2513, |
| 3844, | 3887, | 3510, | 2539, | 2107, | 3186, | 3943, | 2925, | 3242, |
| 3863, | 4028, | 4010, | 3719, | 2858, | 3833, | 3007, | 1850, | 2792, |
| 2982, | 18840, | 2995, | 3579, | 3835, | 3854, | 3508, | 1635, | 24797, |
| 2773, | 5769, | 3634, | 29167, | 5530, | 9000, | 8750, | 1972, | 4983, |
| 8333, | 3667, | 3166, | 3271, | 2241, | 1792, | 2666, | 6478, | 3808, |
| 3729, | 4120, | 6300, | 14987, | 570, | 2600, | 2733, | 3859, | 6825, |
| 3708, | 5314, | 2366, | 2066, | 3767, | 7859, | 4283, | 1700, | 4768, |
| 3083, | 2667, | 1647, | 3400, | 16000, | 2875, | 5041, | 6958, | 5509, |
| 9699, | 3621, | 4709, | 3015, | 2292, | 2360, | 2623, | 3972, | 3522, |
| 6858, | 8334, | 2868, | 3418, | 8667, | 2283, | 5817, | 5119, | 5316, |
| 7603, | 3791, | 3132, | 8550, | 2269, | 4009, | 4158, | 9200, | 5849, |
| 3000, | 2583, | 6000, | 5417, | 3036, | 4006, | 12841, | 3200, | 1853, |
| 1299, | 4950, | 3596, | 4887, | 7660, | 5955, | 3365, | 3717, | 9560, |
| 2799, | 4226, | 1442, | 3750, | 3167, | 4692, | 2275, | 1828, | 3748, |
| 3600, | 1800, | 3941, | 4695, | 3410, | 5649, | 5821, | 2645, | 1928, |
| 3086, | 4230, | 4616, | 11500, | 2132, | 3366, | 8080, | 3357, | 3029, |
| 2609, | 4945, | 5726, | 10750, | 7100, | 4300, | 3208, | 1875, | 4755, |
| 5266, | 1000, | 3846, | 2395, | 1378, | 3988, | 8566, | 5695, | 2958, |
| 3273, | 4133, | 6782, | 2484, | 1977, | 4188, | 1759, | 4288, | 4843, |
| 13650, | 4652, | 3816, | 3052, | 11417, | 7333, | 3800, | 2071, | 2929, |
| 3572, | 7451, | 5050, | 14583, | 2214, | 5568, | 10408, | 2137, | 2957, |
| 3692, | 23803, | 3865, | 10513, | 6080, | 20166, | 2718, | 3459, | 4895, |
| 3316, | 14999, | 4200, | 5042, | 6950, | 2698, | 11757, | 2330, | 14866, |
| 1538, | 4860, | 6277, | 2577, | 9166, | 2281, | 3254, | 39999, | 9538, |
| 2980, | 1863, | 7933, | 9323, | 3707, | 2439, | 2237, | 1820, | 51763, |
| 4344, | 3497, | 2045, | 5516, | 6400, | 1916, | 4600, | 33846, | 3625, |
| 39147, | 2178, | 2383, | 674, | 9328, | 4885, | 12000, | 6033, | 3858, |
| 4191, | 1907, | 3416, | 11000, | 4923, | 3992, | 3917, | 4408, | 3244, |
| 3975, | 2479, | 3430, | 7787, | 5703, | 3173, | 3850, | 150, | 3727, |
| 2221, | 2971, | 7578, | 4735, | 4758, | 2491, | 3716, | 3155, | 5500, |
| 5746, | 3463, | 3812, | 3315, | 5819, | 2510, | 2965, | 3406, | 6050, |
| 9703, | 6608, | 2882, | 1809, | 1668, | 3427, | 2661, | 16250, | 6045, |
| 5250, | 14683, | 4931, | 6083, | 2060, | 3481, | 7200, | 5166, | 4095, |
| 4708, | 2876, | 3237, | 11146, | 2833, | 2620, | 3993, | 3103, | 4100, |


```
sns.countplot(y = 'Gender', hue = 'Loan_Status', data = df)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f99cc032c90>
```



```
df['Gender'].value_counts().plot(kind='pie', autopct='%1.2f%%', figsize=(6, 6))
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f99e0c59f50>
```



The percentage of males who applied for a loan is greater than the one of females.

```
df['Loan_Status'].value_counts().plot(kind='pie', autopct='%1.2f%%', figsize=(6, 6))
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f99cd11d410>



According to the pie chart, there are more approved loans that disapproved.

```
grid=sns.FacetGrid(df, row='Gender', col='Married', size=2.2, aspect=1.6)
grid.map(plt.hist, 'ApplicantIncome', alpha=.5, bins=10)
grid.add_legend()
```


```
/usr/local/lib/python3.7/dist-packages/seaborn/axisgrid.py:316: UserWarning: The `size`
warnings.warn(msg, UserWarning)
<seaborn.axisgrid.FacetGrid at 0x7f99cb69e8d0>
```



Males have the highest income according to the data. Males that are married have greater income than unmarried male. And the same goes for female.

```
grid=sns.FacetGrid(df, row= 'gender', col= 'education', size=(2,2), aspect=1.0)
grid.map(plt.hist, 'ApplicantIncome', alpha=.5, bins=10)
grid.add_legend()
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/axisgrid.py:316: UserWarning: The `size`
warnings.warn(msg, UserWarning)
<seaborn.axisgrid.FacetGrid at 0x7f99cb69e3d0>
```



A graduate who is a male has more income than a one without and the same goes for females.

Here I am exploring the distribution of the numerical variables mainly the Applicant income and the Loan amount.

What can be noticed are quite a few outliers.

```
sns.distplot(df.ApplicantIncome,kde=False)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: `di
warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7f99c262d090>
```

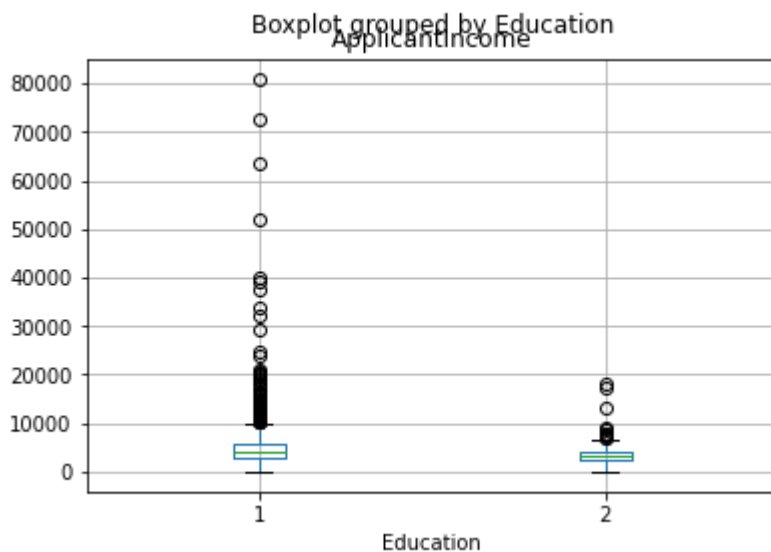


People with better education should normally have a higher income, we can check that by plotting the education level against the income.

```
150 | |
```

```
df.boxplot(column='ApplicantIncome', by = 'Education')
```

```
/usr/local/lib/python3.7/dist-packages/numpy/core/_asarray.py:83: VisibleDeprecationWarn
return array(a, dtype, copy=False, order=order)
<matplotlib.axes._subplots.AxesSubplot at 0x7f99c2587fd0>
```



We can conclude that there is no substantial different between the mean income of graduate and non-graduates. However, there are a higher number of graduates with very high incomes, which are appearing to be the outliers.

```
sns.boxplot(x='Education',y='ApplicantIncome',data=df)
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f99c24be3d0>




The distributions shows that the graduates have more outliers which means that the people with huge income are most likely to be educated.



```
temp3 = pd.crosstab(df['Credit_History'], df['Loan_Status'])
temp3.plot(kind='bar', stacked=True, color=['blue', 'green'], grid=False)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f99c23e5490>



This shows that the chances of getting a loan are higher if the applicant has a valid credit history.

▼ Data processing:

Checking if there are certain missing values that need to be fixed.

```
total = df.isnull().sum().sort_values(ascending=False)
percent = (df.isnull().sum()/df.isnull().count()).sort_values(ascending=False)
missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
missing_data.head(20)
```

| | Total | Percent |
|--------------------------|-------|----------|
| Loan_Status | 367 | 0.374108 |
| Credit_History | 79 | 0.080530 |
| Self_Employed | 55 | 0.056065 |
| LoanAmount | 27 | 0.027523 |
| Dependents | 25 | 0.025484 |
| Gender | 24 | 0.024465 |
| Loan_Amount_Term | 20 | 0.020387 |
| Married | 3 | 0.003058 |
| Property_Area | 0 | 0.000000 |
| CoapplicantIncome | 0 | 0.000000 |
| ApplicantIncome | 0 | 0.000000 |

Filling the missing values, for categorical we can fill them with the mode (the value with the highest frequency). The best practice is to use mode with data points such as salary field or any other kind of money.

```
df['Gender'] = df['Gender'].fillna(
df['Gender'].dropna().mode().values[0] )
df['Married'] = df['Married'].fillna(
df['Married'].dropna().mode().values[0] )
df['Dependents'] = df['Dependents'].fillna(
df['Dependents'].dropna().mode().values[0] )
df['Self_Employed'] = df['Self_Employed'].fillna(
df['Self_Employed'].dropna().mode().values[0] )
df['LoanAmount'] = df['LoanAmount'].fillna(
df['LoanAmount'].dropna().median() )
df['Loan_Amount_Term'] = df['Loan_Amount_Term'].fillna(
df['Loan_Amount_Term'].dropna().mode().values[0] )
df['Credit_History'] = df['Credit_History'].fillna(
df['Credit_History'].dropna().mode().values[0] )
```

Checking if there any empty values.

```
df.isnull().all()
```

```
Loan_ID      False
Gender       False
Married      False
Dependents   False
Education    False
```

```
Self_Employed      False
ApplicantIncome     False
CoapplicantIncome   False
LoanAmount          False
Loan_Amount_Term    False
Credit_History     False
Property_Area       False
Loan_Status         False
dtype: bool
```

```
df.isnull().sum()
```

```
Loan_ID            0
Gender             0
Married            0
Dependents         0
Education          0
Self_Employed     0
ApplicantIncome    0
CoapplicantIncome  0
LoanAmount         0
Loan_Amount_Term   0
Credit_History    0
Property_Area      0
Loan_Status        367
dtype: int64
```

Some people might have a low income, but strong CoapplicantIncome, so a good idea would be to combine them in a TotalIncome column.

```
df['LoanAmount_log']=np.log(df['LoanAmount'])
df['TotalIncome']= df['ApplicantIncome'] + df['CoapplicantIncome']
df['TotalIncome_log']=np.log(df['TotalIncome'])
```

```
sns.distplot(df.TotalIncome,kde=False)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: `di
warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7f99c2305950>
```

▼ Modeling:

Encoding to numeric data in order to start the training of the models.

```
#drop the uniques loan id
```

```
df.drop('Loan_ID', axis = 1, inplace = True)
```

```
df['Gender'].value_counts()
```

```
2.0    799
1.0    182
Name: Gender, dtype: int64
```

```
df['Dependents'].value_counts()
```

```
0      570
1      160
2      160
3+       91
Name: Dependents, dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 981 entries, 0 to 613
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                981 non-null   float64
1   Married               981 non-null   float64
2   Dependents            981 non-null   object
3   Education             981 non-null   int64
4   Self_Employed         981 non-null   float64
5   ApplicantIncome       981 non-null   int64
6   CoapplicantIncome     981 non-null   float64
7   LoanAmount            981 non-null   float64
8   Loan_Amount_Term      981 non-null   float64
9   Credit_History        981 non-null   float64
10  Property_Area         981 non-null   int64
11  Loan_Status           614 non-null   float64
12  LoanAmount_log        981 non-null   float64
13  TotalIncome           981 non-null   float64
14  TotalIncome_log       981 non-null   float64
dtypes: float64(11), int64(3), object(1)
memory usage: 142.6+ KB
```

Need to convert the object values to numeric ones - Dependents needs to become an int.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 981 entries, 0 to 613
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Gender                 981 non-null    float64
1   Married                981 non-null    float64
2   Dependents             981 non-null    object
3   Education              981 non-null    int64
4   Self_Employed          981 non-null    float64
5   ApplicantIncome         981 non-null    int64
6   CoapplicantIncome       981 non-null    float64
7   LoanAmount              981 non-null    float64
8   Loan_Amount_Term        981 non-null    float64
9   Credit_History          981 non-null    float64
10  Property_Area           981 non-null    int64
11  Loan_Status             614 non-null    float64
12  LoanAmount_log          981 non-null    float64
13  TotalIncome             981 non-null    float64
14  TotalIncome_log         981 non-null    float64
dtypes: float64(11), int64(3), object(1)
memory usage: 142.6+ KB
```

Heatmaps are very useful to find relations between two variables in a dataset and this way the user gets a visualisation of the numeric data. No correlations are extremely high. Each square shows the correlation between the variables on each axis.

- The correlations between LoanAmount and ApplicantIncome can be explained:

The close to 1 the correlation is the more positively correlated they are; that is as one increases so does the other and the closer to 1 the stronger this relationship is. It is noticeable that the correlation between the ApplicantIncome and LoanAmount is 0.57, which mean that they have a positive correlation, but not strong.

```
sns.heatmap(df.corr(), annot = True, cmap = 'magma')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f99c2277e10>



Importing sklearn libraries

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import f1_score
```

Splitting into train and test set after choosing the right features X and labels y

```
y = df['Loan_Status']
```

```
X = df.drop('Loan_Status', axis = 1)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=0)
```

```
numeric_features = df.select_dtypes(include=[np.number])
```

```
numeric_features.columns
```

```
Index(['Gender', 'Married', 'Education', 'Self_Employed', 'ApplicantIncome',
       'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Credit_History',
       'Property_Area', 'Loan_Status', 'LoanAmount_log', 'TotalIncome',
       'TotalIncome_log'],
      dtype='object')
```

```
# use only those input features with numeric data type
```

```
df = df.select_dtypes(include=["int64", "float64"])
```

```
# set the target and predictors
```

```
y = df.Loan_Status # target
```

```
# use only those input features with numeric data type
```

```
df_temp = df.select_dtypes(include=["int64", "float64"])
```

```
X = df_temp.drop(["Loan_Status"],axis=1) # predictors
```

Logistic Regression

```
model = LogisticRegression()
model.fit(X_train, y_train)
y_reg=model.predict(X_test)
evaluation = f1_score(y_test, y_reg)
evaluation
```

ValueError Traceback (most recent call last)

<ipython-input-129-ae0c7bcd60e1> in <module>()

```
1 model = LogisticRegression()
----> 2 model.fit(X_train, y_train)
3 y_reg=model.predict(X_test)
4 evaluation = f1_score(y_test, y_reg)
5 evaluation
```

5 frames

/usr/local/lib/python3.7/dist-packages/numpy/core/_asarray.py in asarray(a, dtype, order)

```
81
82     """
---> 83     return array(a, dtype, copy=False, order=order)
84
85
```

ValueError: could not convert string to float: '3+'

SEARCH STACK OVERFLOW

Decision tree:

1. Creating classifier
2. Fitting classifier with train data

```
tree = DecisionTreeClassifier()
tree.fit(X_train, y_train)
```

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',
                        max_depth=None, max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort='deprecated',
                        random_state=None, splitter='best')
```

Do predictions on a test set. **Testing** the model by testing the test data.

```
y_tree=tree.predict(X_test)
print(y_tree)
```

```
[0 1 1 0 1 1 1 1 0 1 0 1 1 1 1 1 1 1 0 0 0 1 0 0 0 1 0 1 0 0 1 0 1 1 1 0 1
 1 1 1 0 0 1 0 0 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 0 0 1 0 0 1 1 1 0 1 1 1 1 0
 1 1 0 1 1 0 1 1 1 1 1 0 1 0 1 1 1 1 0 0 0 0 1 1 1 1 0 1 0 1 1 1 1 1 1 0 0
 1 1 1 0 0 0 1 1 0 1 0 0]
```

Evaluate classssifier, measure accuracy, which is 0.76

```
evaluation = f1_score(y_test, y_tree)
evaluation
```

```
0.7619047619047619
```

Random forests

```
forest = RandomForestClassifier()
forest.fit(X_train, y_train)
```

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='gini', max_depth=None, max_features='auto',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=100,
                        n_jobs=None, oob_score=False, random_state=None,
                        verbose=0, warm_start=False)
```

Testing the model by testing the test data.

```
y_forest=forest.predict(X_test)
print(y_forest)
```

```
[1 1 1 1 1 0 1 0 0 1 1 1 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 0 0 1 1 1 1 1 0 1
 1 1 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1
 1 1 1 0 1 0 0 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1
 0 1 1 0 0 1 1 1 1 1 0 1]
```

Result of the **accuracy**.

```
evaluation_f= f1_score(y_test, y_forest)
evaluation_f
```

```
0.8795811518324607
```


Conclusion

From the Exploratory Data Analysis, it can be concluded:

1. There amount of male applicants seems to be greater than the female ones and they tend to live in the semisuburban areas.
2. There are more positive than negative loan statuses - more approvals.
3. The distributions shows that the graduates have more outliers which means that the people with huge income are most likely to be educated.
4. Males have the highest income according to the data. Males that are married have greater income than unmarried male. And the same goes for female. Therefore, there is a greater chance for educated and married people to receive a loan than applicant who are not.

From the Modelling, it can be concluded:

1. After the exploring of different types of modelling, that the more accurate model is Random forest than Decision tree.
2. From the evaluation of the three models, it can be noticed that the Logistic Regression performed better than others